

1.Uber Ride

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import warnings
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib.cbook import boxplot_stats
import plotly.express as px
from geopy import Point, distance
from math import *
warnings.filterwarnings("ignore")
```

```
df = pd.read_csv("uber.csv")
```

```
df.head()
```

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.00000003	7.5	
1	27835199	2009-07-17 20:04:56.00000002	7.7	
2	44984355	2009-08-24 21:45:00.000000061	12.9	
3	25894730	2009-06-26 08:22:21.00000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	

	pickup_datetime	pickup_longitude	pickup_latitude	\
0	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	2014-08-28 17:47:00 UTC	-73.925023	40.744085	

	dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1
1	-73.994710	40.750325	1
2	-73.962565	40.772647	1
3	-73.965316	40.803349	3
4	-73.973082	40.761247	5

```
df = df.drop(["Unnamed: 0", "key"], axis=1)
```

```
df.head()
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	\
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	

2	12.9	2009-08-24 21:45:00 UTC	-74.005043
40.740770			
3	5.3	2009-06-26 08:22:21 UTC	-73.976124
40.790844			
4	16.0	2014-08-28 17:47:00 UTC	-73.925023
40.744085			

	dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1
1	-73.994710	40.750325	1
2	-73.962565	40.772647	1
3	-73.965316	40.803349	3
4	-73.973082	40.761247	5

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 193786 entries, 0 to 199999
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	fare_amount	193786 non-null	float64
1	pickup_datetime	193786 non-null	datetime64[ns, UTC]
2	pickup_longitude	193786 non-null	float64
3	pickup_latitude	193786 non-null	float64
4	dropoff_longitude	193786 non-null	float64
5	dropoff_latitude	193786 non-null	float64
6	passenger_count	193786 non-null	int64
7	distance_km	193786 non-null	float64
8	pickup_hr	193786 non-null	int64
9	day	193786 non-null	int64
10	month	193786 non-null	int64
11	year	193786 non-null	int64
12	day_of_week	193786 non-null	int64
13	day_name	193786 non-null	object

```
dtypes: datetime64[ns, UTC](1), float64(6), int64(6), object(1)
memory usage: 22.2+ MB
```

```
df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"],
errors="coerce")
```

```
df.describe().T
```

	count	mean	std	min
25% \				
fare_amount	200000.0	11.359955	9.901776	-52.000000
6.000000				
pickup_longitude	200000.0	-72.527638	11.437787	-1340.648410
73.992065				
pickup_latitude	200000.0	39.935885	7.720539	-74.015515

```

40.734796
dropoff_longitude 199999.0 -72.525292 13.117408 -3356.666300 -
73.991407
dropoff_latitude 199999.0 39.923890 6.794829 -881.985513
40.733823
passenger_count 200000.0 1.684535 1.385997 0.000000
1.000000

```

	50%	75%	max
fare_amount	8.500000	12.500000	499.000000
pickup_longitude	-73.981823	-73.967154	57.418457
pickup_latitude	40.752592	40.767158	1644.421482
dropoff_longitude	-73.980093	-73.963658	1153.572603
dropoff_latitude	40.753042	40.768001	872.697628
passenger_count	1.000000	2.000000	208.000000

```
df.isna().sum()
```

```

fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 1
dropoff_latitude  1
passenger_count  0
dtype: int64

```

```
df = df.dropna()
```

```

def distance_transform(longitude1, latitude1, longitude2, latitude2):
    distance = []
    for pos in range(len(longitude1)):
        long1,lat1,long2,lati2 = map(radians,
[longitude1[pos],latitude1[pos],longitude2[pos],latitude2[pos]])
        dist_long = long2 - long1
        dist_lati = lati2 - lat1
        a = sin(dist_lati/2)**2 + cos(lat1) * cos(lati2) *
sin(dist_long/2)**2
        c = 2 * asin(sqrt(a))*6371
        distance.append(c)
    return distance

```

```

df["distance_km"] =
distance_transform(df["pickup_longitude"].to_numpy(),
df["pickup_latitude"].to_numpy(),
df["dropoff_longitude"].to_numpy(), df["dropoff_latitude"].to_numpy())
df = df.assign(pickup_hr = df.pickup_datetime.dt.hour,
               day= df.pickup_datetime.dt.day,
               month = df.pickup_datetime.dt.month,

```

```

year = df.pickup_datetime.dt.year,
day_of_week = df.pickup_datetime.dt.dayofweek,
day_name=df.pickup_datetime.dt.day_name())

```

```
df.head()
```

	fare_amount		pickup_datetime	pickup_longitude
pickup_latitude \				
0	7.5	2015-05-07	19:52:06+00:00	-73.999817
40.738354				
1	7.7	2009-07-17	20:04:56+00:00	-73.994355
40.728225				
2	12.9	2009-08-24	21:45:00+00:00	-74.005043
40.740770				
3	5.3	2009-06-26	08:22:21+00:00	-73.976124
40.790844				
4	16.0	2014-08-28	17:47:00+00:00	-73.925023
40.744085				

	dropoff_longitude	dropoff_latitude	passenger_count	
distance_km \				
0	-73.999512	40.723217	1	1.683323
1	-73.994710	40.750325	1	2.457590
2	-73.962565	40.772647	1	5.036377
3	-73.965316	40.803349	3	1.661683
4	-73.973082	40.761247	5	4.475450

	pickup_hr	day	month	year	day_of_week	day_name
0	19	7	5	2015	3	Thursday
1	20	17	7	2009	4	Friday
2	21	24	8	2009	0	Monday
3	8	26	6	2009	4	Friday
4	17	28	8	2014	3	Thursday

```

def find_outliers(df):
    q1 = df.quantile(0.25)
    q3 = df.quantile(0.75)
    IQR = q3-q1
    outliers = df[((df<(q1-1.5*IQR)) | (df>(q3+1.5*IQR)))]
    return outliers

```

```

outliers = find_outliers(df['fare_amount'])
print('number of outliers:' + str(len(outliers)))
print('max outlier value:' + str(outliers.max()))
print('min outlier value:' + str(outliers.min()))
outliers

```

```
number of outliers:17166
max outlier value:499.0
min outlier value:-52.0
```

```
6          24.50
30          25.70
34          39.50
39          29.00
48          56.80
```

```
...
199976     49.70
199977     43.50
199982     57.33
199985     24.00
199997     30.90
```

```
Name: fare_amount, Length: 17166, dtype: float64
```

```
outliers = find_outliers(df['passenger_count'])
print('number of outliers:' + str(len(outliers)))
print('max outlier value:' + str(outliers.max()))
print('min outlier value:' + str(outliers.min()))
outliers
```

```
number of outliers:22557
max outlier value:208
min outlier value:4
```

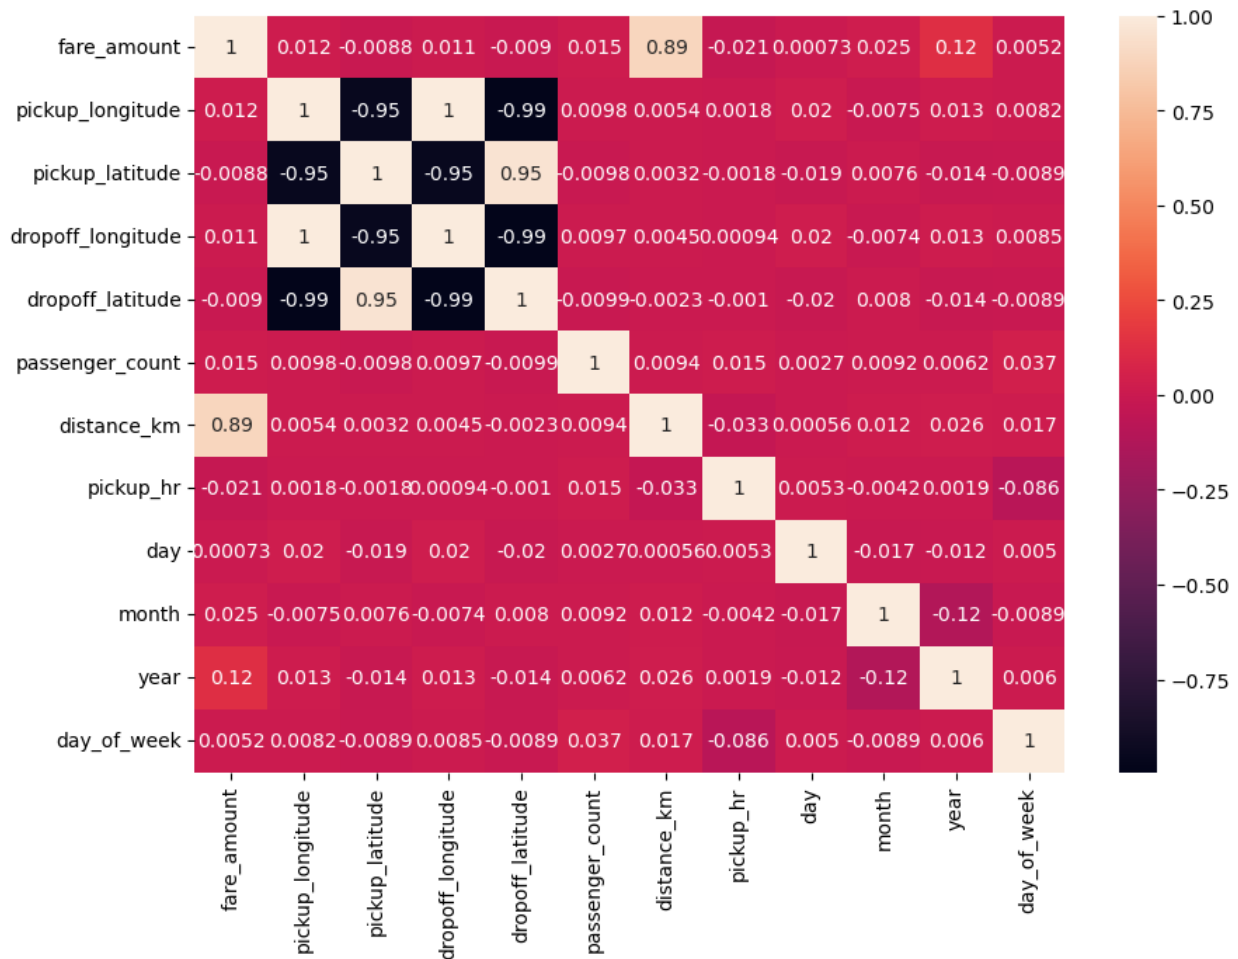
```
4          5
6          5
12         5
24         5
29         5
```

```
..
199958     5
199959     5
199962     4
199969     5
199985     5
```

```
Name: passenger_count, Length: 22557, dtype: int64
```

```
df.drop(df[df['distance_km'] == 0].index, inplace = True)
df.drop(df[df['distance_km'] > 60].index, inplace = True)
df.drop(df[df['fare_amount'] > 100].index, inplace = True)
df.drop(df[df['fare_amount'] < 0].index, inplace = True)
df.drop(df[df['passenger_count'] > 6].index, inplace = True)
```

```
plt.figure(figsize=(10,7))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
x = df[["year", "distance_km"]]
y = df["fare_amount"]

scaler = StandardScaler()
x = scaler.fit_transform(x)


x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.3, random_state=42)

model = LinearRegression()
model.fit(x_train, y_train)

LinearRegression()

y_pred = model.predict(x_test)

sns.regplot(x=y_test, y=y_pred, color="red", line_kws={"color" :
"blue"})
plt.show()
```



```
print(f"Mean absolute error {metrics.mean_absolute_error(y_test,  
y_pred)}")  
print(f"Mean squared error {metrics.mean_squared_error(y_test,  
y_pred)}")  
print(f"Root mean squared error  
{np.sqrt(metrics.mean_squared_error(y_test, y_pred))}")
```


```
Mean absolute error 2.243879503459476  
Mean squared error 18.32381910645457  
Root mean squared error 4.280633026370582
```

```
model = RandomForestRegressor()  
model.fit(x_train, y_train)
```

```
RandomForestRegressor()
```

```
y_pred = model.predict(x_test)
```

```
sns.regplot(x=y_test, y=y_pred, color="red", line_kws={"color" :  
"blue"})  
plt.show()
```



```
print(f"Mean absolute error {metrics.mean_absolute_error(y_test,
y_pred)}")
print(f"Mean squared error {metrics.mean_squared_error(y_test,
y_pred)}")
print(f"Root mean squared error
{np.sqrt(metrics.mean_squared_error(y_test, y_pred))}")
```

```
Mean absolute error 2.500337233122361
Mean squared error 21.283639254402114
Root mean squared error 4.613419475226821
```

```
def read_data(path: str) -> pd.DataFrame:
    """
    Read data from csv file.

    Args:
        path (str): path to csv file.

    Returns:
        pd.DataFrame: dataframe of csv file.
    """
    df = pd.read_csv(path)

    return df
```


2.Email Classification

```
def basic_info(df: pd.DataFrame) -> pd.DataFrame:
    """
    Get basic information of dataframe.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        pd.DataFrame: dataframe of basic information.
    """
    return df.info()

def distance_transform(longitude1: np.ndarray, latitude1: np.ndarray,
    longitude2: np.ndarray, latitude2: np.ndarray) -> list:
    """
    Calculate distance between two points.

    Args:
        longitude1 (np.ndarray): array of longitude of first point.
        latitude1 (np.ndarray): array of latitude of first point.
        longitude2 (np.ndarray): array of longitude of second point.
        latitude2 (np.ndarray): array of latitude of second point.

    Returns:
        list: list of distance between two points.
    """
    distance = []
    for pos in range(len(longitude1)):
        long1,lati1,long2,lati2 = map(radians,
        [longitude1[pos],latitude1[pos],longitude2[pos],latitude2[pos]])
        dist_long = long2 - long1
        dist_lati = lati2 - lati1
        a = sin(dist_lati/2)**2 + cos(lati1) * cos(lati2) *
sin(dist_long/2)**2
        c = 2 * asin(sqrt(a))*6371
        distance.append(c)

    return distance

def find_outliers(df: pd.DataFrame) -> pd.DataFrame:
    """
    Find outliers in dataframe.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        pd.DataFrame: dataframe of outliers.
    """
```

```

q1 = df.quantile(0.25)
q3 = df.quantile(0.75)
IQR = q3-q1
outliers = df[((df<(q1-1.5*IQR)) | (df>(q3+1.5*IQR)))]

return outliers

def preprocess(df: pd.DataFrame) -> pd.DataFrame:
    """
    Preprocess dataframe.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        pd.DataFrame: dataframe after preprocessing.
    """
    df = df.drop(["Unnamed: 0", "key"], axis=1)
    df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"],
errors="coerce")
    df = df.dropna()
    df["distance_km"] =
distance_transform(df["pickup_longitude"].to_numpy(),
df["pickup_latitude"].to_numpy(),
df["dropoff_longitude"].to_numpy(), df["dropoff_latitude"].to_numpy())
    df = df.assign(pickup_hr = df.pickup_datetime.dt.hour,
                    day= df.pickup_datetime.dt.day,
                    month = df.pickup_datetime.dt.month,
                    year = df.pickup_datetime.dt.year,
                    day_of_week = df.pickup_datetime.dt.dayofweek,
                    day_name=df.pickup_datetime.dt.day_name())
    outliers = find_outliers(df['fare_amount'])
    print('number of outliers for fare amount:' + str(len(outliers)))
    print('max outlier value for fare amount:' + str(outliers.max()))
    print('min outlier value for fare amount:' + str(outliers.min()))
    print(outliers)
    outliers = find_outliers(df['passenger_count'])
    print('number of outliers for fare amount:' + str(len(outliers)))
    print('max outlier value for fare amount:' + str(outliers.max()))
    print('min outlier value for fare amount:' + str(outliers.min()))
    print(outliers)
    df.drop(df[df['distance_km'] == 0].index, inplace = True)
    df.drop(df[df['distance_km'] > 60].index, inplace = True)
    df.drop(df[df['fare_amount'] > 100].index, inplace = True)
    df.drop(df[df['fare_amount'] < 0].index, inplace = True)
    df.drop(df[df['passenger_count'] > 6].index, inplace = True)

    return df

```

```

def visualize_correlation(df: pd.DataFrame) -> None:
    """
    Visualize correlation between features.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        None.
    """
    plt.figure(figsize=(10,7))
    sns.heatmap(df.corr(), annot=True)
    plt.show()

def split_data(df: pd.DataFrame) -> tuple:
    """
    Split data into train and test set.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        tuple: tuple of train and test set.
    """
    x = df[["year", "distance_km"]]
    y = df["fare_amount"]
    scaler = StandardScaler()
    scaler.fit_transform(x)
    x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.3, random_state=42)

    return x_train, x_test, y_train, y_test

def create_model(model_name: str) -> object:
    """
    Create model.

    Args:
        model_name (str): name of model.

    Returns:
        object: model.
    """
    if model_name == "LR":
        model = LinearRegression()
    elif model_name == "RFR":
        model = RandomForestRegressor()

    return model

```

```

def train_model(model: object, x_train: np.ndarray, y_train:
np.ndarray) -> None:
    """
    Train model.

    Args:
        model (object): model.
        x_train (np.ndarray): array of train set.
        y_train (np.ndarray): array of train set.

    Returns:
        None.
    """
    model.fit(x_train, y_train)

def test_model(model: object, x_test: np.ndarray) -> np.ndarray:
    """
    Test model.

    Args:
        model (object): model.
        x_test (np.ndarray): array of test set.

    Returns:
        np.ndarray: array of predicted value.
    """
    y_pred = model.predict(x_test)

    return y_pred

def reg_line(y_test: np.ndarray, y_pred: np.ndarray) -> None:
    """
    Visualize regression line.

    Args:
        y_test (np.ndarray): test value.
        y_pred (np.ndarray): predicted value.
    """
    sns.regplot(x=y_test, y=y_pred, color="red", line_kws={"color" :
"blue"})
    plt.show()

def metrics_model(y_test: np.ndarray, y_pred: np.ndarray) -> None:
    """
    Calculate metrics of model.

    Args:
        y_test (np.ndarray): test value.
        y_pred (np.ndarray): predicted value.

```

```

Returns:
    None.
"""
    print(f"Mean absolute error {metrics.mean_absolute_error(y_test,
y_pred)}")
    print(f"Mean squared error {metrics.mean_squared_error(y_test,
y_pred)}")
    print(f"Root mean squared error
{np.sqrt(metrics.mean_squared_error(y_test, y_pred))}")

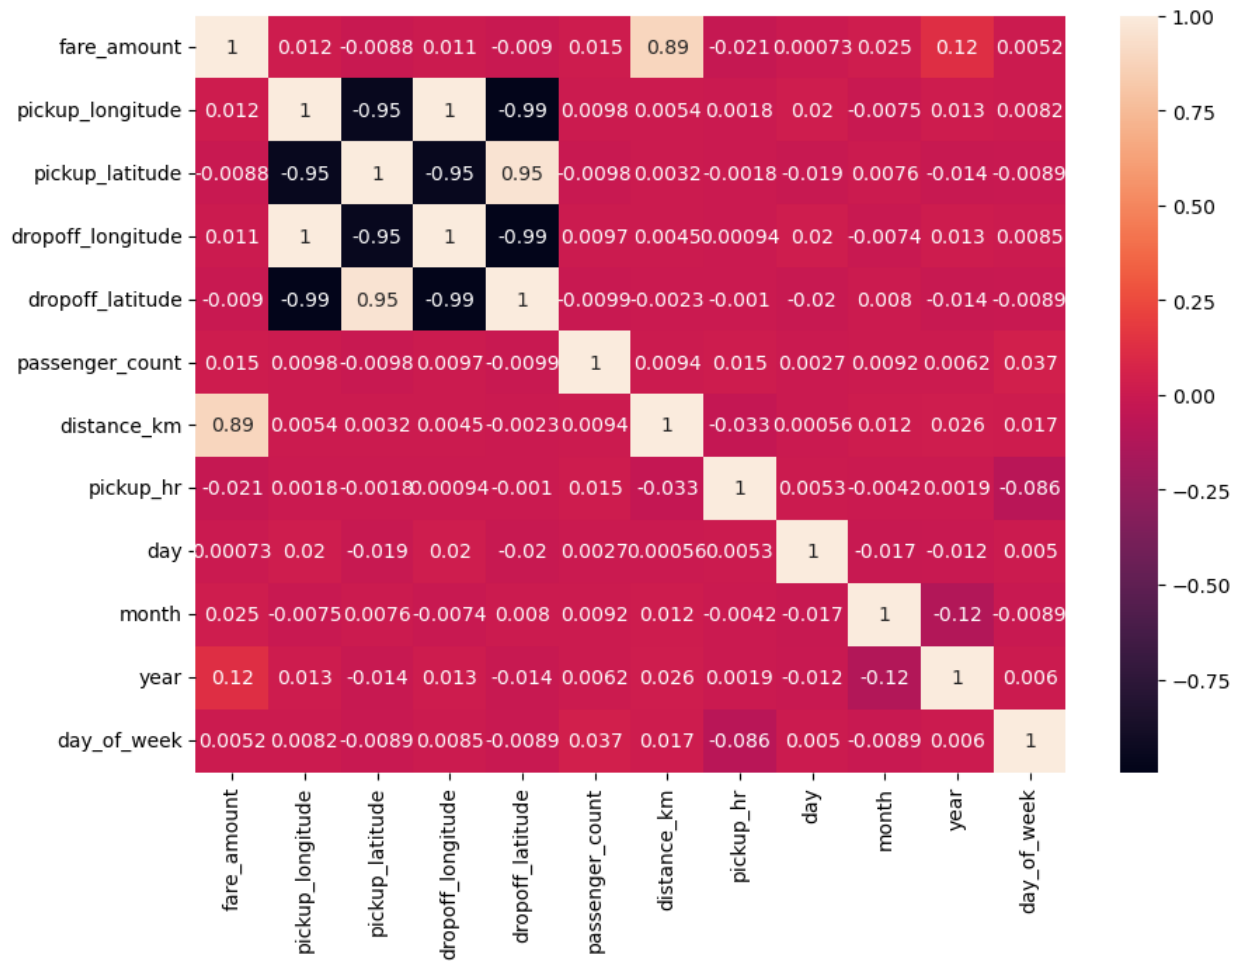
df = read_data("/kaggle/input/uber-fares-dataset/uber.csv")
print(basic_info(df))
df = preprocess(df)
print("\nCorrelation Matrix:\n")
visualize_correlation(df)
x_train, x_test, y_train, y_test = split_data(df)
model = create_model("LR")
train_model(model, x_train, y_train)
y_pred = test_model(model, x_test)
print("\nRegression Line:\n")
reg_line(y_test, y_pred)
print("\nModel Metrics:\n")
metrics_model(y_test, y_pred)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
None
number of outliers for fare amount:17166
max outlier value for fare amount:499.0
min outlier value for fare amount:-52.0
6         24.50
30        25.70
34        39.50
39        29.00
48        56.80
...


```

```
199976    49.70
199977    43.50
199982    57.33
199985    24.00
199997    30.90
Name: fare_amount, Length: 17166, dtype: float64
number of outliers for fare amount:22557
max outlier value for fare amount:208
min outlier value for fare amount:4
4         5
6         5
12        5
24        5
29        5
..
199958    5
199959    5
199962    4
199969    5
199985    5
Name: passenger_count, Length: 22557, dtype: int64

Correlation Matrix:
```



Regression Line:



Model Metrics:

Mean absolute error 2.2438795034594645

Mean squared error 18.323819106454575

Root mean squared error 4.280633026370583

3.Neutral Network Classifier

Import libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.utils import resample
from sklearn import metrics
from tqdm.notebook import tqdm
%matplotlib inline
warnings.filterwarnings("ignore")

df = pd.read_csv("emails.csv")

df.head()

df.shape

df.describe().T
```

Without upsampling

```
df = df.drop("Email No.", axis=1)

df.isna().sum()

sns.distplot(x=df["Prediction"])
plt.show()

x = df.drop("Prediction", axis=1)
y = df[["Prediction"]]

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2)
```

KNN with elbow plot

```
k_values = [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
accuracy_values = []

for i in tqdm(range(len(k_values))):
    model = KNeighborsClassifier(n_neighbors=k_values[i])
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    accuracy_values.append(accuracy)
```

```

accuracy_values
px.line(x=k_values, y=accuracy_values)

optimal_k = -1
optimal_accuracy = -1
for i in list(zip(k_values, accuracy_values)):
    if i[1] > optimal_accuracy:
        optimal_k = i[0]
        optimal_accuracy = i[1]

knn_model = KNeighborsClassifier(n_neighbors=optimal_k)
knn_model.fit(x_train, y_train)
y_pred = knn_model.predict(x_test)
print(metrics.classification_report(y_test, y_pred))

```

SVM

```

svm_model = SVC()
svm_model.fit(x_train, y_train)
y_pred = svm_model.predict(x_test)
print(metrics.classification_report(y_test, y_pred))

```

With upsampling

```

spam_data = df[df["Prediction"] == 1]
ham_data = df[df["Prediction"] == 0]

spam_upsample = resample(spam_data,
                        replace=True,
                        n_samples=int(0.8*len(ham_data)),
                        random_state=42)

new_df = ham_data
new_df = new_df.append(spam_upsample)

new_df.head()

new_df.shape

new_df = new_df.sample(frac=1)

sns.distplot(new_df["Prediction"])
plt.show()

x = new_df.drop("Prediction", axis=1)
y = new_df[["Prediction"]]

```

```
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2)
```

KNN with elbow plot

```
k_values = [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
accuracy_values = []

for i in tqdm(range(len(k_values))):
    model = KNeighborsClassifier(n_neighbors=k_values[i])
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    accuracy_values.append(accuracy)

px.line(x=k_values, y=accuracy_values)

optimal_k = -1
optimal_accuracy = -1
for i in list(zip(k_values, accuracy_values)):
    if i[1] > optimal_accuracy:
        optimal_k = i[0]
        optimal_accuracy = i[1]

knn_model = KNeighborsClassifier(n_neighbors=optimal_k)
knn_model.fit(x_train, y_train)
y_pred = knn_model.predict(x_test)
print(metrics.classification_report(y_test, y_pred))
```

SVM

```
svm_model = SVC()
svm_model.fit(x_train, y_train)
y_pred = svm_model.predict(x_test)
print(metrics.classification_report(y_test, y_pred))
```

Functions

```
def read_data(path: str) -> pd.DataFrame:
    """
    Read data from csv file.
```

```

    Args:
        path (str): path to csv file.

    Returns:
        pd.DataFrame: dataframe of csv file.
    """
    df = pd.read_csv(path)
    return df

def basic_info(df: pd.DataFrame) -> pd.DataFrame:
    """
    Get basic information of dataframe.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        pd.DataFrame: dataframe of basic information.
    """
    return df.info()

def preprocess(df: pd.DataFrame) -> pd.DataFrame:
    df = df.drop("Email No.", axis=1)
    return df

def split_data(df: pd.DataFrame) -> tuple:
    """
    Split data into train and test set.

    Args:
        df (pd.DataFrame): dataframe.

    Returns:
        tuple: tuple of train and test set.
    """
    x = df.drop("Prediction", axis=1)
    y = df[["Prediction"]]
    x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.3, random_state=42)
    return x_train, x_test, y_train, y_test

def knn_model_with_elbow_method(x_train: np.ndarray, x_test:
np.ndarray, y_train: np.ndarray, y_test: np.ndarray, k_values: list) -
> np.ndarray:
    """
    KNN model with elbow method.

    Args:
        x_train (np.ndarray): x_train data.
        x_test (np.ndarray): x_test data.

```

```

    y_train (np.ndarray): y_train data.
    y_test (np.ndarray): y_test data.
    k_values (list): list of k values.

Returns:
    np.ndarray: y_pred data.
"""
accuracy_values = []
for i in tqdm(range(len(k_values))):
    model = KNeighborsClassifier(n_neighbors=k_values[i])
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    accuracy_values.append(accuracy)
fig = px.line(x=k_values, y=accuracy_values, title="K value vs
Accuracy")
fig.update_layout(xaxis_title="K values", yaxis_title="Accuracy
values")
fig.show()
optimal_k = -1
optimal_accuracy = -1
for i in list(zip(k_values, accuracy_values)):
    if i[1] > optimal_accuracy:
        optimal_k = i[0]
        optimal_accuracy = i[1]
knn_model = KNeighborsClassifier(n_neighbors=optimal_k)
knn_model.fit(x_train, y_train)
y_pred = knn_model.predict(x_test)
return y_pred

def svm_model(x_train: np.ndarray, x_test: np.ndarray, y_train:
np.ndarray, y_test: np.ndarray) -> np.ndarray:
    """
    SVM model.

    Args:
        x_train (np.ndarray): x_train data.
        x_test (np.ndarray): x_test data.
        y_train (np.ndarray): y_train data.
        y_test (np.ndarray): y_test data.

    Returns:
        np.ndarray: y_pred data.
    """
    svm_model = SVC()
    svm_model.fit(x_train, y_train)
    y_pred = svm_model.predict(x_test)
    return y_pred

def metrics_report(y_test: np.ndarray, y_pred: np.ndarray) -> None:

```

```

"""
    Print metrics report.

    Args:
        y_test (np.ndarray): y_test data.
        y_pred (np.ndarray): y_pred data.
"""
print(metrics.classification_report(y_test, y_pred))

def upsample_data(df: pd.DataFrame) -> pd.DataFrame:
    """
        Upsample data.

        Args:
            df (pd.DataFrame): dataframe.

        Returns:
            pd.DataFrame: upsampled dataframe.
    """
    spam_data = df[df["Prediction"] == 1]
    ham_data = df[df["Prediction"] == 0]
    spam_upsample = resample(
        spam_data,
        replace=True,
        n_samples=int(0.8*len(ham_data)),
        random_state=42
    )
    new_df = ham_data
    new_df = new_df.append(spam_upsample)
    new_df = new_df.sample(frac=1)
    return new_df

df =
read_data("/kaggle/input/email-spam-classification-dataset-csv/emails.
csv")
basic_info(df)
df = preprocess(df)
df = upsample_data(df)
x_train, x_test, y_train, y_test = split_data(df)
k_values = [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
y_pred_knn = knn_model_with_elbow_method(x_train, x_test, y_train,
y_test, k_values)
y_pred_svm = svm_model(x_train, x_test, y_train, y_test)
print("Metrics for KNN-\n")
metrics_report(y_test, y_pred_knn)
print("Metrics for SVM-\n")
metrics_report(y_test, y_pred_svm)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5172 entries, 0 to 5171

```

Columns: 3002 entries, Email No. to Prediction
dtypes: int64(3001), object(1)
memory usage: 118.5+ MB

```
{"model_id": "47e76ce383a84e9bb256a446d794f3f0", "version_major": 2, "version_minor": 0}
```

Metrics for KNN-

	precision	recall	f1-score	support
0	0.96	0.87	0.91	1128
1	0.85	0.96	0.90	855
accuracy			0.91	1983
macro avg	0.90	0.91	0.91	1983
weighted avg	0.91	0.91	0.91	1983

Metrics for SVM-

	precision	recall	f1-score	support
0	0.79	0.92	0.85	1128
1	0.86	0.67	0.76	855
accuracy			0.81	1983
macro avg	0.82	0.80	0.80	1983
weighted avg	0.82	0.81	0.81	1983

4.Gradient Descent Algorithm

Implement Gradient Descent Algorithm to find the local minima of a function. For example, find the local minima of the function $y=(x+3)**2$ starting from the point $x=2$.

```
import matplotlib.pyplot as plt

def cost_function(x):
    # ithe given function yenar
    return (x + 3) ** 2

def gradient(x):
    # ithe derivate of given function yenar
    return 2 * (x + 3)

learning_rate = 0.1
initial_x = 2.0
num_iterations = 100

x_values = []
y_values = []
x = initial_x
for i in range(num_iterations):
    x_values.append(x)
    y_values.append(cost_function(x))
    gradient_value = gradient(x)
    x = x - learning_rate * gradient_value

    print(f'Iteration {i+1}: x = {x}, Cost = {cost_function(x)}')

print(f'Optimal x: {x}')
```


Iteration 1: x = 1.0, Cost = 16.0
Iteration 2: x = 0.19999999999999996, Cost = 10.240000000000002
Iteration 3: x = -0.44000000000000017, Cost = 6.553599999999998
Iteration 4: x = -0.9520000000000001, Cost = 4.194304
Iteration 5: x = -1.3616000000000001, Cost = 2.6843545599999996
Iteration 6: x = -1.6892800000000001, Cost = 1.7179869183999996
Iteration 7: x = -1.951424, Cost = 1.099511627776
Iteration 8: x = -2.1611392, Cost = 0.7036874417766399
Iteration 9: x = -2.32891136, Cost = 0.4503599627370493
Iteration 10: x = -2.463129088, Cost = 0.28823037615171165
Iteration 11: x = -2.5705032704, Cost = 0.1844674407370954
Iteration 12: x = -2.6564026163200003, Cost = 0.11805916207174093
Iteration 13: x = -2.725122093056, Cost = 0.07555786372591429
Iteration 14: x = -2.7800976744448, Cost = 0.04835703278458515
Iteration 15: x = -2.82407813955584, Cost = 0.030948500982134555
Iteration 16: x = -2.8592625116446717, Cost = 0.019807040628566166
Iteration 17: x = -2.8874100093157375, Cost = 0.012676506002282305
Iteration 18: x = -2.90992800745259, Cost = 0.008112963841460692
Iteration 19: x = -2.927942405962072, Cost = 0.005192296858534868
Iteration 20: x = -2.9423539247696575, Cost = 0.0033230699894623056

Iteration 21: x = -2.953883139815726, Cost = 0.002126764793255884
Iteration 22: x = -2.9631065118525806, Cost = 0.0013611294676837786
Iteration 23: x = -2.9704852094820646, Cost = 0.0008711228593176078
Iteration 24: x = -2.9763881675856516, Cost = 0.0005575186299632732
Iteration 25: x = -2.981110534068521, Cost = 0.00035681192317650156
Iteration 26: x = -2.984888427254817, Cost = 0.00022835963083295564
Iteration 27: x = -2.9879107418038537, Cost = 0.00014615016373308945
Iteration 28: x = -2.990328593443083, Cost = 9.353610478917726e-05
Iteration 29: x = -2.9922628747544664, Cost = 5.986310706507345e-05
Iteration 30: x = -2.993810299803573, Cost = 3.83123885216492e-05
Iteration 31: x = -2.995048239842858, Cost = 2.451992865385725e-05
Iteration 32: x = -2.9960385918742864, Cost = 1.5692754338469342e-05
Iteration 33: x = -2.9968308734994293, Cost = 1.0043362776619253e-05
Iteration 34: x = -2.9974646987995435, Cost = 6.427752177036323e-06
Iteration 35: x = -2.997971759039635, Cost = 4.113761393302886e-06
Iteration 36: x = -2.998377407231708, Cost = 2.6328072917135587e-06
Iteration 37: x = -2.998701925785366, Cost = 1.684996666971388e-06
Iteration 38: x = -2.998961540628293, Cost = 1.0783978666865378e-06
Iteration 39: x = -2.9991692325026342, Cost = 6.901746346793842e-07
Iteration 40: x = -2.9993353860021075, Cost = 4.417117661946878e-07
Iteration 41: x = -2.999468308801686, Cost = 2.826955303647891e-07
Iteration 42: x = -2.9995746470413485, Cost = 1.8092513943361614e-07
Iteration 43: x = -2.9996597176330786, Cost = 1.1579208923763523e-07
Iteration 44: x = -2.999727774106463, Cost = 7.410693711203819e-08
Iteration 45: x = -2.99978221928517, Cost = 4.7428439751781807e-08
Iteration 46: x = -2.9998257754281363, Cost = 3.035420144107846e-08
Iteration 47: x = -2.999860620342509, Cost = 1.9426688922339734e-08
Iteration 48: x = -2.999888496274007, Cost = 1.243308091029743e-08
Iteration 49: x = -2.9999107970192056, Cost = 7.9571717826062e-09
Iteration 50: x = -2.9999286376153647, Cost = 5.092589940842615e-09
Iteration 51: x = -2.9999429100922916, Cost = 3.259257562149415e-09
Iteration 52: x = -2.999954328073833, Cost = 2.0859248397837384e-09
Iteration 53: x = -2.9999634624590668, Cost = 1.3349918974486118e-09
Iteration 54: x = -2.9999707699672533, Cost = 8.543948143723039e-10
Iteration 55: x = -2.999976615973803, Cost = 5.468126811899669e-10
Iteration 56: x = -2.9999812927790424, Cost = 3.499601159582557e-10
Iteration 57: x = -2.9999850342232337, Cost = 2.2397447421860056e-10
Iteration 58: x = -2.999988027378587, Cost = 1.433436634977776e-10
Iteration 59: x = -2.9999904219028695, Cost = 9.173994464198049e-11
Iteration 60: x = -2.9999923375222957, Cost = 5.871356456950638e-11
Iteration 61: x = -2.9999938700178364, Cost = 3.757668132666189e-11
Iteration 62: x = -2.999995096014269, Cost = 2.4049076048192486e-11
Iteration 63: x = -2.9999960768114153, Cost = 1.5391408670843192e-11
Iteration 64: x = -2.9999968614491324, Cost = 9.850501548782124e-12
Iteration 65: x = -2.999997489159306, Cost = 6.3043209907745444e-12
Iteration 66: x = -2.9999979913274446, Cost = 4.034765434809332e-12
Iteration 67: x = -2.9999983930619556, Cost = 2.5822498785634223e-12
Iteration 68: x = -2.9999987144495646, Cost = 1.6526399220522305e-12
Iteration 69: x = -2.9999989715596516, Cost = 1.0576895502961154e-12

```
Iteration 70: x = -2.9999991772477212, Cost = 6.769213121895138e-13
Iteration 71: x = -2.999999341798177, Cost = 4.3322963956744853e-13
Iteration 72: x = -2.9999994734385416, Cost = 2.7726696951023927e-13
Iteration 73: x = -2.9999995787508333, Cost = 1.7745086041172427e-13
Iteration 74: x = -2.9999996630006667, Cost = 1.1356855066350352e-13
Iteration 75: x = -2.9999997304005332, Cost = 7.268387247253274e-14
Iteration 76: x = -2.9999997843204267, Cost = 4.651767834410857e-14
Iteration 77: x = -2.9999998274563415, Cost = 2.977131407892966e-14
Iteration 78: x = -2.9999998619650734, Cost = 1.9053640961475125e-14
Iteration 79: x = -2.9999998895720585, Cost = 1.2194330254575965e-14
Iteration 80: x = -2.999999911657647, Cost = 7.804371331543109e-15
Iteration 81: x = -2.9999999293261177, Cost = 4.994797639633387e-15
Iteration 82: x = -2.999999943460894, Cost = 3.1966704893653676e-15
Iteration 83: x = -2.9999999547687155, Cost = 2.045869097124455e-15
Iteration 84: x = -2.9999999638149726, Cost = 1.309356209304147e-15
Iteration 85: x = -2.9999999710519782, Cost = 8.379879636702507e-16
Iteration 86: x = -2.9999999768415826, Cost = 5.363122967489605e-16
Iteration 87: x = -2.999999981473266, Cost = 3.432398699193347e-16
Iteration 88: x = -2.9999999851786128, Cost = 2.1967351938118145e-16
Iteration 89: x = -2.99999998814289, Cost = 1.4059105661644777e-16
Iteration 90: x = -2.999999990514312, Cost = 8.997827286453327e-17
Iteration 91: x = -2.9999999924114498, Cost = 5.758609463330129e-17
Iteration 92: x = -2.9999999939291597, Cost = 3.685510164371068e-17
Iteration 93: x = -2.9999999951433276, Cost = 2.358726677741145e-17
Iteration 94: x = -2.999999996114662, Cost = 1.5095850047368678e-17
Iteration 95: x = -2.99999999689173, Cost = 9.661342926036557e-18
Iteration 96: x = -2.9999999975133838, Cost = 6.18326035608692e-18
Iteration 97: x = -2.999999998010707, Cost = 3.957287334634505e-18
Iteration 98: x = -2.9999999984085655, Cost = 2.532663894166083e-18
Iteration 99: x = -2.999999998726852, Cost = 1.6209053445792253e-18
Iteration 100: x = -2.999999998981482, Cost = 1.0373792396055266e-18
Optimal x: -2.99999998981482
```

```
plt.plot(x_values, y_values, 'ro-')
plt.title('Gradient Descent Visualization for  $y = (x + 3)^2$  by AB')
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```

Gradient Descent Visualization for $y = (x + 3)^2$ by AB



5.K-Nearest NEighbors

Import libraries

```
import pandas as pd
import numpy as np
import plotly.express as px
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.utils import resample
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from mlxtend.plotting import plot_confusion_matrix
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

Data loading and preprocessing

```
df = pd.read_csv("diabetes.csv")
```

df

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	Pedigree	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 768 entries, 0 to 767

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	Pedigree	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

df.describe().T

	count	mean	std	min	25%
50% \					
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000
3.0000					
Glucose	768.0	120.894531	31.972618	0.000	99.00000
117.0000					
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000
72.0000					
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000
23.0000					
Insulin	768.0	79.799479	115.244002	0.000	0.00000
30.5000					
BMI	768.0	31.992578	7.884160	0.000	27.30000

32.0000					
Pedigree	768.0	0.471876	0.331329	0.078	0.24375
0.3725					
Age	768.0	33.240885	11.760232	21.000	24.00000
29.0000					
Outcome	768.0	0.348958	0.476951	0.000	0.00000
0.0000					


	75%	max
Pregnancies	6.00000	17.00
Glucose	140.25000	199.00
BloodPressure	80.00000	122.00
SkinThickness	32.00000	99.00
Insulin	127.25000	846.00
BMI	36.60000	67.10
Pedigree	0.62625	2.42
Age	41.00000	81.00
Outcome	1.00000	1.00

```
df["Outcome"].value_counts()
```

```
0    500
1    268
```

```
Name: Outcome, dtype: int64
```

```
sns.countplot(data=df, x=df["Outcome"])
plt.show()
```



Upsampling

```
negative_data = df[df["Outcome"] == 0]
positive_data = df[df["Outcome"] == 1]


positive_upsample = resample(positive_data,
                             replace=True,
                             n_samples=int(0.9*len(negative_data)),
                             random_state=42)

new_df = negative_data
new_df = new_df.append(positive_upsample)

new_df.shape
(950, 9)

new_df = new_df.sample(frac=1)

sns.countplot(data=new_df, x=new_df["Outcome"])
plt.show()
```



```
x = new_df.drop("Outcome", axis=1)
y = new_df[["Outcome"]]

scaler = MinMaxScaler()
scaled_values = scaler.fit_transform(x)

x_train, x_test, y_train, y_test = train_test_split(scaled_values, y,
test_size=0.2)
```

KNN with elbow plot

```
k_values = [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29, 31,
33, 35, 37, 39, 41, 43, 45, 47, 49]
accuracy_values = []

for i in tqdm(range(len(k_values))):
    model = KNeighborsClassifier(n_neighbors=k_values[i])
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    accuracy_values.append(accuracy)
```



```
{"model_id": "bacfa33a6b4e4ce8833eaf0ffdb45be7", "version_major": 2, "version_minor": 0}
```

```
px.line(x=k_values, y=accuracy_values)
```

```
optimal_k = -1
```

```
optimal_accuracy = -1
```

```
for i in list(zip(k_values, accuracy_values)):
```

```
    if i[1] > optimal_accuracy:
```

```
        optimal_k = i[0]
```

```
        optimal_accuracy = i[1]
```

```
knn_model = KNeighborsClassifier(n_neighbors=optimal_k)
```

```
knn_model.fit(x_train, y_train)
```

```
KNeighborsClassifier(n_neighbors=1)
```

```
y_pred = knn_model.predict(x_test)
```


```
print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.78	0.86	102
1	0.79	0.95	0.87	88
accuracy			0.86	190
macro avg	0.87	0.87	0.86	190
weighted avg	0.88	0.86	0.86	190

```
cm = metrics.confusion_matrix(y_test, y_pred)
```

```
plot_confusion_matrix(cm)
```

```
plt.show()
```



```
y_score = model.predict_proba(x_test)[: ,1]


false_positive_rate, true_positive_rate, threshold =
metrics.roc_curve(y_test, y_score)

print('roc_auc_score for DecisionTree: ',
metrics.roc_auc_score(y_test, y_score))

roc_auc_score for DecisionTree: 0.7575200534759358

plt.subplots(1, figsize=(10,7))
plt.title('Receiver Operating Characteristic - KNN')
plt.plot(false_positive_rate, true_positive_rate)
plt.plot([0, 1], ls="--")
plt.plot([0, 0], [1, 0] , c=".7"), plt.plot([1, 1] , c=".7")
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Receiver Operating Characteristic - KNN



```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import warnings
from sklearn.preprocessing import StandardScaler
warnings.filterwarnings('ignore')
```

```
df = pd.read_csv("sales_data_sample.csv", encoding="latin")
```

```
df.head()
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	
SALES \					
0	10107	30	95.70	2	2871.00
1	10121	34	81.35	5	2765.90
2	10134	41	94.74	2	3884.34
3	10145	45	83.26	6	3746.70
4	10159	49	100.00	14	5205.27

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	\
0	2/24/2003 0:00	Shipped	1	2	2003	...	
1	5/7/2003 0:00	Shipped	2	5	2003	...	
2	7/1/2003 0:00	Shipped	3	7	2003	...	
3	8/25/2003 0:00	Shipped	3	8	2003	...	
4	10/10/2003 0:00	Shipped	4	10	2003	...	

	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME
DEALSIZE					
0	10022	USA	NaN	Yu	Kwai
Small					
1	51100	France	EMEA	Henriot	Paul
Small					
2	75508	France	EMEA	Da Cunha	Daniel
Medium					
3	90003	USA	NaN	Young	Julie
Medium					
4	NaN	USA	NaN	Brown	Julie
Medium					

```
[5 rows x 25 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2823 entries, 0 to 2822
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
0	ORDERNUMBER	2823 non-null	int64
1	QUANTITYORDERED	2823 non-null	int64
2	PRICEEACH	2823 non-null	float64
3	ORDERLINENUMBER	2823 non-null	int64
4	SALES	2823 non-null	float64
5	ORDERDATE	2823 non-null	object
6	STATUS	2823 non-null	object
7	QTR_ID	2823 non-null	int64
8	MONTH_ID	2823 non-null	int64
9	YEAR_ID	2823 non-null	int64
10	PRODUCTLINE	2823 non-null	object
11	MSRP	2823 non-null	int64
12	PRODUCTCODE	2823 non-null	object
13	CUSTOMERNAME	2823 non-null	object
14	PHONE	2823 non-null	object
15	ADDRESSLINE1	2823 non-null	object
16	ADDRESSLINE2	302 non-null	object
17	CITY	2823 non-null	object
18	STATE	1337 non-null	object
19	POSTALCODE	2747 non-null	object
20	COUNTRY	2823 non-null	object
21	TERRITORY	1749 non-null	object
22	CONTACTLASTNAME	2823 non-null	object
23	CONTACTFIRSTNAME	2823 non-null	object
24	DEALSIZE	2823 non-null	object

```
dtypes: float64(2), int64(7), object(16)
```

```
memory usage: 551.5+ KB
```

```
df = df[['ORDERLINENUMBER', 'SALES']]
```

```
scaler = StandardScaler()
```

```
scaled_values = scaler.fit_transform(df.values)
```

```
wcss = []
```

```
for i in range(1, 11):
```


```
    model = KMeans(n_clusters=i, init='k-means++')
```

```
    model.fit_predict(scaled_values)
```

```
    wcss.append(model.inertia_)
```

```
plt.plot(range(1, 11), wcss, 'ro-')
```

```
plt.show()
```



```
model = KMeans(n_clusters=7, init='k-means++')
clusters = model.fit_predict(scaled_values)
clusters
```

```
array([3, 3, 0, ..., 4, 3, 6])
```

```
df['cluster'] = clusters
```

```
df
```


	ORDERLINENUMBER	SALES	cluster
0	2	2871.00	3
1	5	2765.90	3
2	2	3884.34	0
3	6	3746.70	0
4	14	5205.27	5
...
2818	15	2244.40	1
2819	1	3978.51	0
2820	4	5417.57	4
2821	1	2116.16	3
2822	9	3079.44	6

```
[2823 rows x 3 columns]
```

```
model.inertia_
```

993.4283577026391

```
plt.scatter(df['ORDERLINENUMBER'], df['SALES'], c=df['cluster'])  
plt.show()
```



**SCTR's Pune Institute of Computer Technology
Dhankawadi, Pune**

A MINI-PROJECT REPORT ON

Titanic Survivor Prediction Using Machine Learning

SUBMITTED BY

Aayush Goyal	41302
Divya khandare	41321
Arnav Firke	41322

**Under the guidance of
Prof. V.S. Gaikwad**



**DEPARTMENT OF COMPUTER ENGINEERING
ACADEMIC YEAR 2024-25**

Contents

1	Title	2
2	Introduction	2
3	Problem Statement	3
4	About Dataset	3
5	Objectives and Scope	3
5.1	Objectives of Project:	3
5.2	Scope of Project:	3
6	Methodological Details	4
7	Modern Engineering Tools Used	5
7.1	Google Colab:	5
7.2	Python:	5
7.3	Libraries/Frameworks:	5
8	Conclusion	6

1 Title

Titanic Survivor Prediction Using Machine Learning

2 Introduction

Predicting Titanic survivors is a significant application of machine learning in predictive analytics. Understanding the factors influencing passenger survival can offer valuable insights into historical events and enhance decision-making processes in safety and emergency response systems. This project focuses on building a machine learning model to predict the likelihood of a Titanic passenger's survival based on various features such as age, gender, class, and fare.

The motivation behind this project arises from the historical significance of the Titanic disaster and the wealth of data available to explore survival patterns. By leveraging advanced machine learning algorithms and Titanic passenger data, we aim to create a predictive model that can accurately determine survival outcomes based on passenger characteristics.

We selected algorithms like Logistic Regression, Random Forest, and Support Vector Machines (SVM) for their ability to capture complex relationships between features and their proven effectiveness in classification tasks. By training these models, we aim to uncover the key factors that contributed to survival and provide interpretable results.

The significance of this project lies in its contribution to understanding historical data through modern machine-learning techniques. This approach not only aids in predictive accuracy but also sheds light on critical variables that impacted survival during one of history's most well-known maritime disasters.

3 Problem Statement

To predict Titanic Survivors based on various given features.

4 About Dataset

The dataset used for Titanic survivor prediction is the well-known Titanic.csv which contains detailed information about the passengers aboard the RMS Titanic. This dataset originates from the records of the Titanic disaster and includes various features that describe the passengers, such as their age, gender, ticket class, fare, and whether or not they survived.

The dataset provides a comprehensive snapshot of the 2,224 passengers on board, with data on both survivors and non-survivors, allowing for in-depth analysis. The data covers essential variables like Pclass (ticket class), Sex, Age, SibSp(number of siblings/spouses aboard), Parch (number of parents/children aboard), Fare, Embarked (port of embarkation), and Survived (the target variable indicating survival).

This historical data allows us to explore the relationships between different features and survival outcomes, offering the potential to build a robust predictive model.

5 Objectives and Scope

5.1 Objectives of Project:

The main objective of the Titanic survivor prediction project is to accurately predict whether a passenger survived or not based on given input variables. By analyzing features such as age, gender, ticket class, and fare, this prediction can provide insights into the factors influencing survival during the Titanic disaster. The model can be used in historical analysis, educational purposes, and to better understand the relationship between passenger attributes and survival outcomes.

5.2 Scope of Project:

1. **Data Collection:** Utilize the Titanic dataset, which contains historical information about passengers aboard the Titanic, including variables such as age, gender, class, fare, and survival status. This data will be used for training and testing the machine learning model.
2. **Feature Selection:** Identify and select key features that have a significant impact on survival. This may include variables like passenger class (Pclass), gender (Sex), age, number of siblings/spouses aboard (SibSp), and fare, among others.
3. **Data Preprocessing:** Prepare the dataset by handling missing values, converting categorical data into numerical form (such as gender and embarked location), and scaling or normalizing features if needed to ensure proper model performance.
4. **Model Training:** Train the predictive model using machine learning algorithms such as Logistic Regression, Random Forest, and Support Vector Machines (SVM). These algorithms will analyze the selected features to predict the survival of each passenger.
5. **Model Evaluation:** Evaluate the performance of the trained model using appropriate metrics like accuracy, precision, recall, F1-score, and confusion matrix. Cross-validation and testing on unseen data will ensure the model's reliability.
6. **Model Deployment:** Deploy the trained model into an application that can take passenger data as input and predict the likelihood of survival. This system can be integrated into educational tools, simulations, or analytical software for further exploration of the Titanic disaster.

6 Methodological Details

This is a methodological breakdown of the project:

1. **Data Collection and Loading:** The Titanic survival prediction project begins with the collection of historical data, specifically the Titanic dataset, which is loaded using **Pandas**. The dataset includes essential features such as the age, gender, and class of the passengers, as well as whether they survived or not. The data was sourced from a `.csv` file stored in a cloud directory and was accessed and manipulated within a Python-based Jupyter notebook.
2. **Data Cleaning and Preprocessing:** Once the dataset is loaded, an initial inspection is carried out to identify missing values. The dataset includes several columns that are either irrelevant or contain excessive missing data, such as **PassengerId**, **Name**, **Ticket**, and **Cabin**, which are subsequently dropped. The remaining columns are retained for analysis.
3. **Exploratory Data Analysis:** EDA is conducted using visualizations created with **Seaborn** and **Matplotlib** to understand the distribution of survival across different features like gender (**Sex**) and passenger class (**Pclass**). The missing values in the **Age** and **Embarked** columns are handled through imputation. The **Age** column is imputed using the mean strategy, while the **Embarked** column is filled with its most frequent value (**mode**).
4. **Encoding Categorical Variables:** To facilitate machine learning model training, the categorical columns such as **Sex** and **Embarked** are converted into numerical formats using **Label Encoding**. This conversion is crucial for allowing algorithms to interpret categorical data during the training process.
5. **Feature Selection and Correlation Analysis:** Before proceeding to model development, a heatmap is generated to visualize the correlation matrix of the selected features. This step helps identify the relationships between different variables and survival, aiding in feature selection.
6. **Model Training:** The dataset is then split into training and testing sets using **train_test_split** from **Scikit-learn**, with 80% of the data used for training and 20% for testing. This allows the model to learn from a portion of the data while preserving another portion for evaluation purposes. The primary machine learning model used in this project is a **Random Forest Classifier** due to its robustness and ability to handle non-linear relationships. The model is trained on the selected features using the training set.
7. **Model Evaluation:** After training, predictions are made on the test set using the trained model. The model's performance is evaluated using metrics such as **accuracy**, **confusion matrix**, and **classification report**, which provide a detailed breakdown of how well the model is performing in terms of precision, recall, and F1-score.

7 Modern Engineering Tools Used

The breakdown of the tools and technologies used in the project:

7.1 Google Colab:

Google Colab is used as the primary platform for developing and running the project. It provides a Jupyter notebook environment with access to free GPU and TPU resources, which is useful for machine learning tasks.

7.2 Python:

Python is the programming language used for coding the project. It is a versatile language with rich libraries for data analysis, machine learning, and visualization.

7.3 Libraries/Frameworks:

- **NumPy:** Used for numerical computations and array manipulation.
- **Pandas:** Used for data manipulation and analysis, including reading and writing data in various formats.
- **Matplotlib:** Used for data visualization, including plotting graphs and charts.
- **scikit-learn:** Used for machine learning tasks such as model building, evaluation, and preprocessing.

Hardware Requirements:

Since Google Colab is a cloud-based platform, no specific hardware requirements are needed from the user's end. However, having a stable internet connection is necessary to access and work on Google Colab.

These libraries provide extensive documentation and resources for data analysis, visualization, and machine learning tasks.

Google Colab Environment:

Google Colab provides a pre-configured environment with the following specifications:

- Python runtime environment
- Access to GPU and TPU resources for accelerated computation
- Integrated code editor with support for Jupyter notebooks
- Ability to install additional libraries using pip or conda
- Integration with Google Drive for data storage and access

8 Conclusion

In conclusion, this project has successfully demonstrated the application of machine learning techniques for predicting the survival of Titanic passengers based on various features such as age, gender, ticket class, and fare. Through comprehensive data preprocessing, feature selection, model training, and evaluation, we have developed a predictive model capable of accurately forecasting passenger survival with solid performance metrics.

Machine learning algorithms like Logistic Regression, Random Forest, and Support Vector Machines (SVM) have proven to be effective in capturing the complex relationships between passenger characteristics and survival outcomes. By leveraging historical Titanic data and advanced machine learning methodologies, we gained valuable insights into the factors influencing survival during the disaster.

The practical applications of this predictive model extend to educational tools, historical analysis, and safety simulations, providing a better understanding of survival factors in maritime disasters. The model's predictions can serve as a foundation for further exploration of survival patterns and emergency preparedness.

Moving forward, opportunities for refinement include incorporating additional features such as cabin location or family size, improving model performance through hyperparameter tuning, and deploying the model into real-world simulations for broader evaluation.

Overall, this project represents a significant contribution toward analyzing historical data using machine learning, and it showcases the potential of data-driven approaches to understanding human survival in critical scenarios like the Titanic disaster.